

SAM: Optimizing Multithreaded Cores for Speculative Parallelism

MALEEN ABEYDEERA, SUVINAY SUBRAMANIAN, MARK JEFFREY,
JOEL EMER, DANIEL SANCHEZ

PACT 2017



Executive Summary

Analyzes the interplay between hardware multithreading and speculative parallelism

(eg: Thread Level Speculation and Transactional Memory)

Conventional multithreading causes performance pathologies on speculative workloads

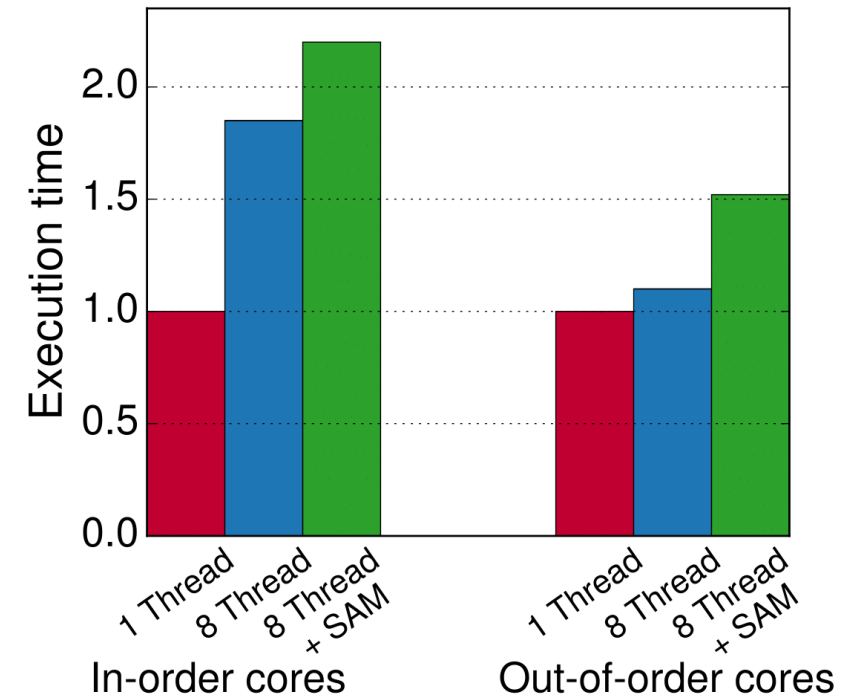
- Increase in aborted work
- Inefficient use of speculation resources

Why? All threads are treated equally

Speculation Aware Multithreading (SAM)

- Prioritize threads running tasks more likely to commit

SAM makes multithreading more useful



Outline

Background on speculative parallelism

Pitfalls of speculative parallelism with conventional multithreading

SAM on in-order cores

SAM on out-of-order cores

Background on Speculative Parallelism

Parallelize tasks when the dependences are not known in advance

Hardware executes all tasks in parallel, aborting upon conflicts

Which task to abort? Conflict resolution policy

Speculative Parallelism



```
graph TD; A[Speculative Parallelism] --> B[Ordered]; A --> C[Unordered];
```

Ordered

e.g. Thread-Level Speculation (TLS)

(Program order dictates the conflict resolution order)

Unordered

e.g. Hardware Transactional Memory

(Any execution order is valid, but high-performance conflict resolution policies define an order)

Implicit order among all tasks in any speculative system

Baseline System - Swarm [Jeffrey, MICRO' 15]

```
void desTask(Timestamp ts , GateInput* input) {  
    Gate* g = input ->gate ();  
    bool toggledOutput = g.simulateToggle(input);  
    if ( toggledOutput ) {  
        for (GateInput* i : g-> connectedInputs ()) {  
            swarm::enqueue(desTask , ts+delay(g,i), i);  
        }  
    }  
}
```

Timestamped tasks

Tasks create children tasks
(function ptr, timestamp, args)

Tasks appear to execute in timestamp order

Unordered execution via equal timestamps

Swarm Microarchitecture

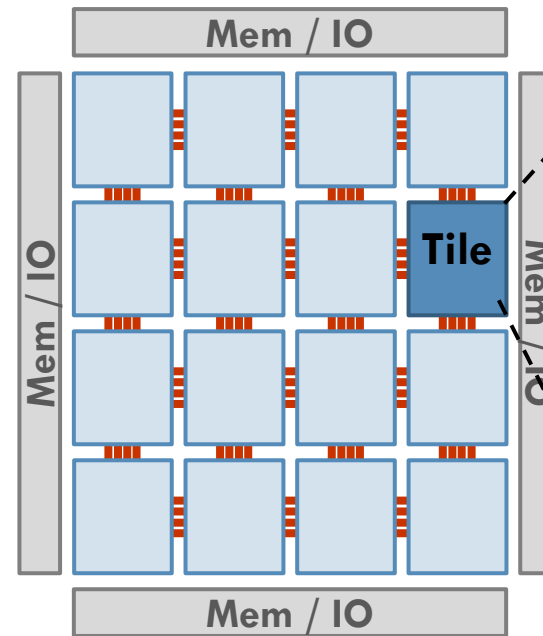
Equal timestamps:
global order via Virtual Time (VT)



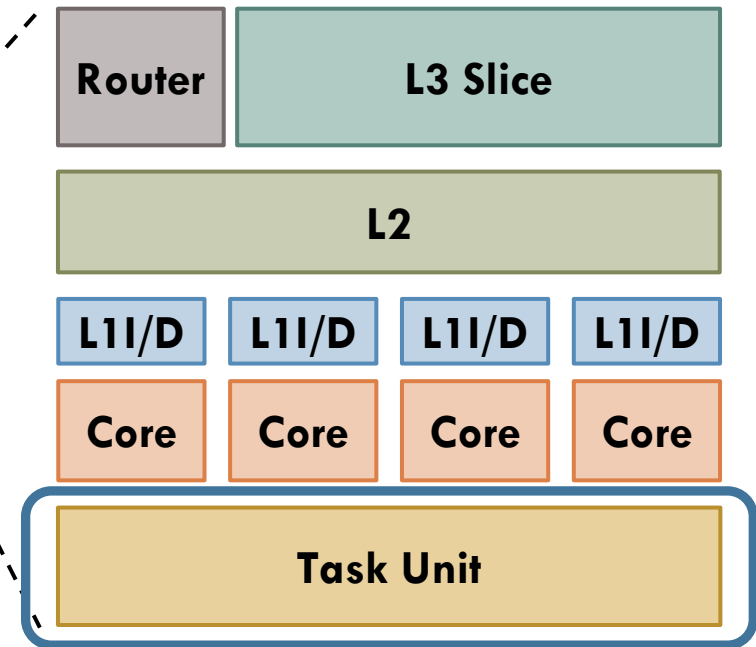
Tasks execute out-of-order,
but commit in VT order

Commit queue: state of tasks waiting to commit

16-tile, 64-core CMP



Tile Organization



Outline

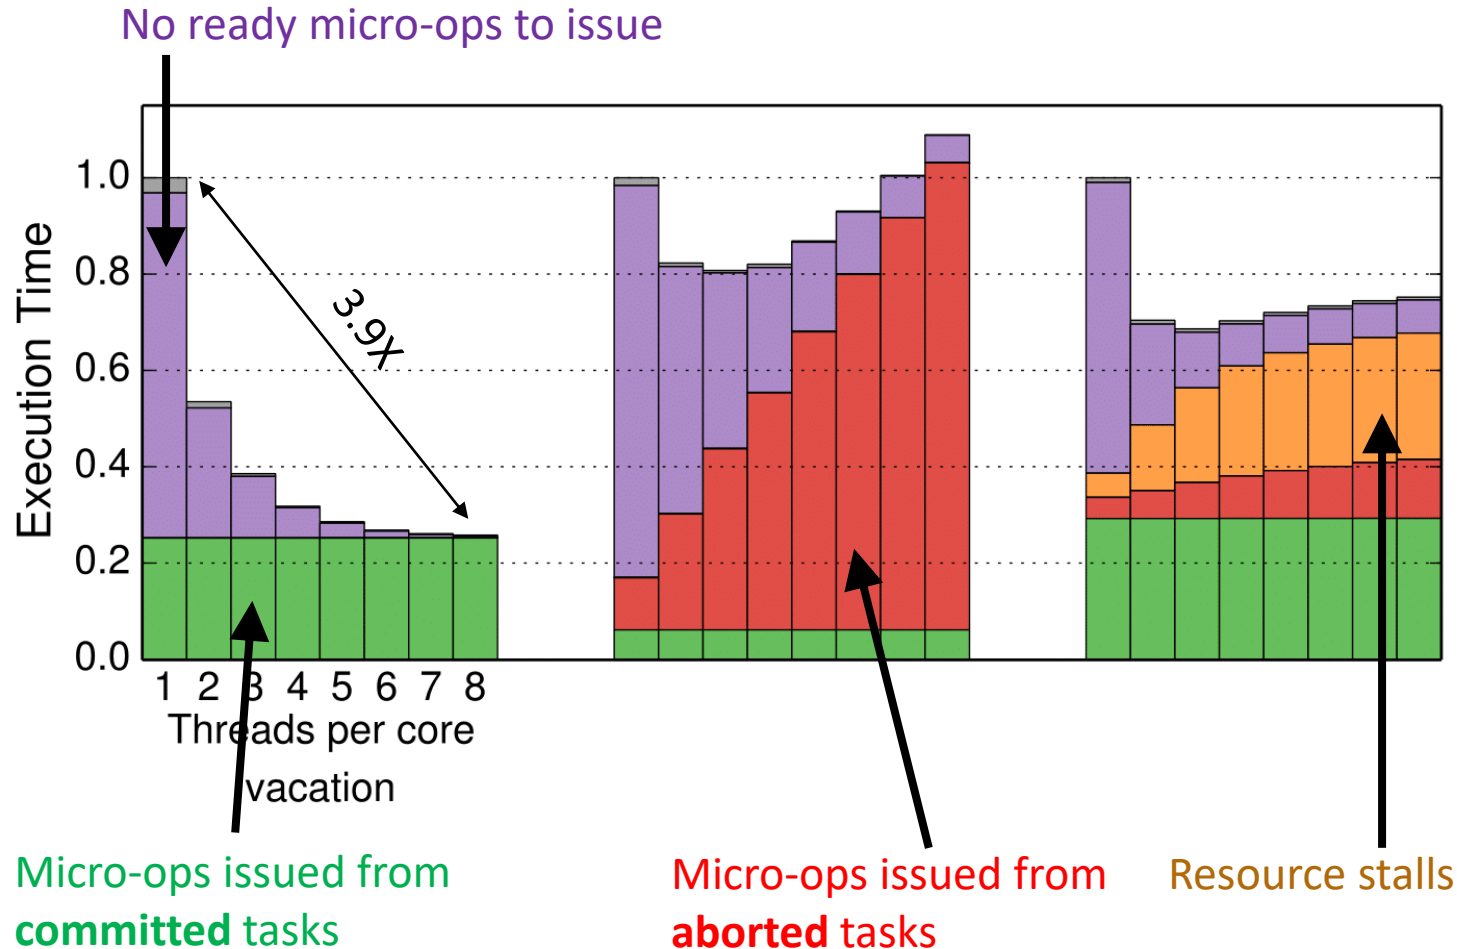
Background on speculative parallelism

Pitfalls of speculative parallelism with conventional multithreading

SAM on in-order cores

SAM on out-of-order cores

Pitfalls of Speculation-Oblivious Multithreading



System configuration:

64-core SMT system

In-order core with 2-wide issue

Speculation-oblivious round-robin order

Insights:

1. Multithreading can be highly beneficial

However, multithreading can also lead to:

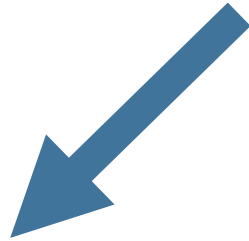
2. Increased aborts

3. Inefficient use of speculation resources

Unlikely-to-commit tasks hurt the throughput of likely-to-commit ones

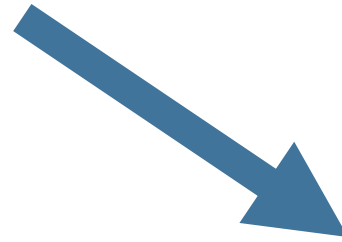
Speculation-Aware Multithreading

**Prioritize threads according to their
conflict resolution priorities**



Reduce Aborts

(focus resources on tasks likely to commit)



Reduce Speculation Resource Stalls

(tasks commit early)

Outline

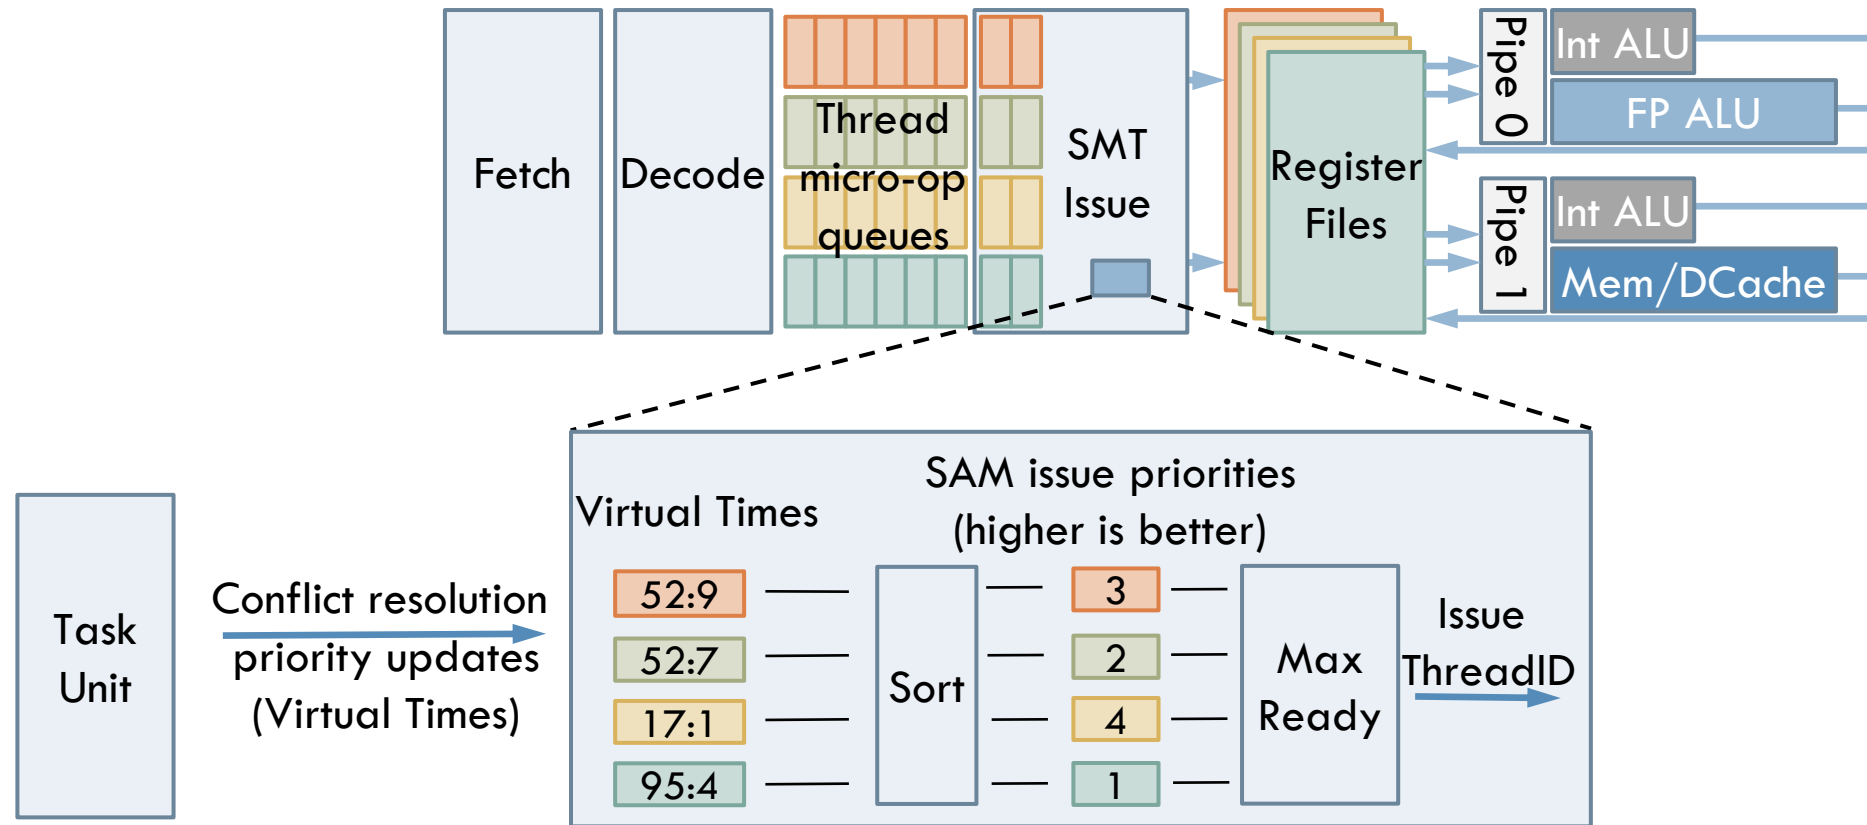
Background on speculative parallelism

Pitfalls of speculative parallelism with conventional multithreading

SAM on in-order cores

SAM on out-of-order cores

SAM on in-order cores



Experimental Methodology

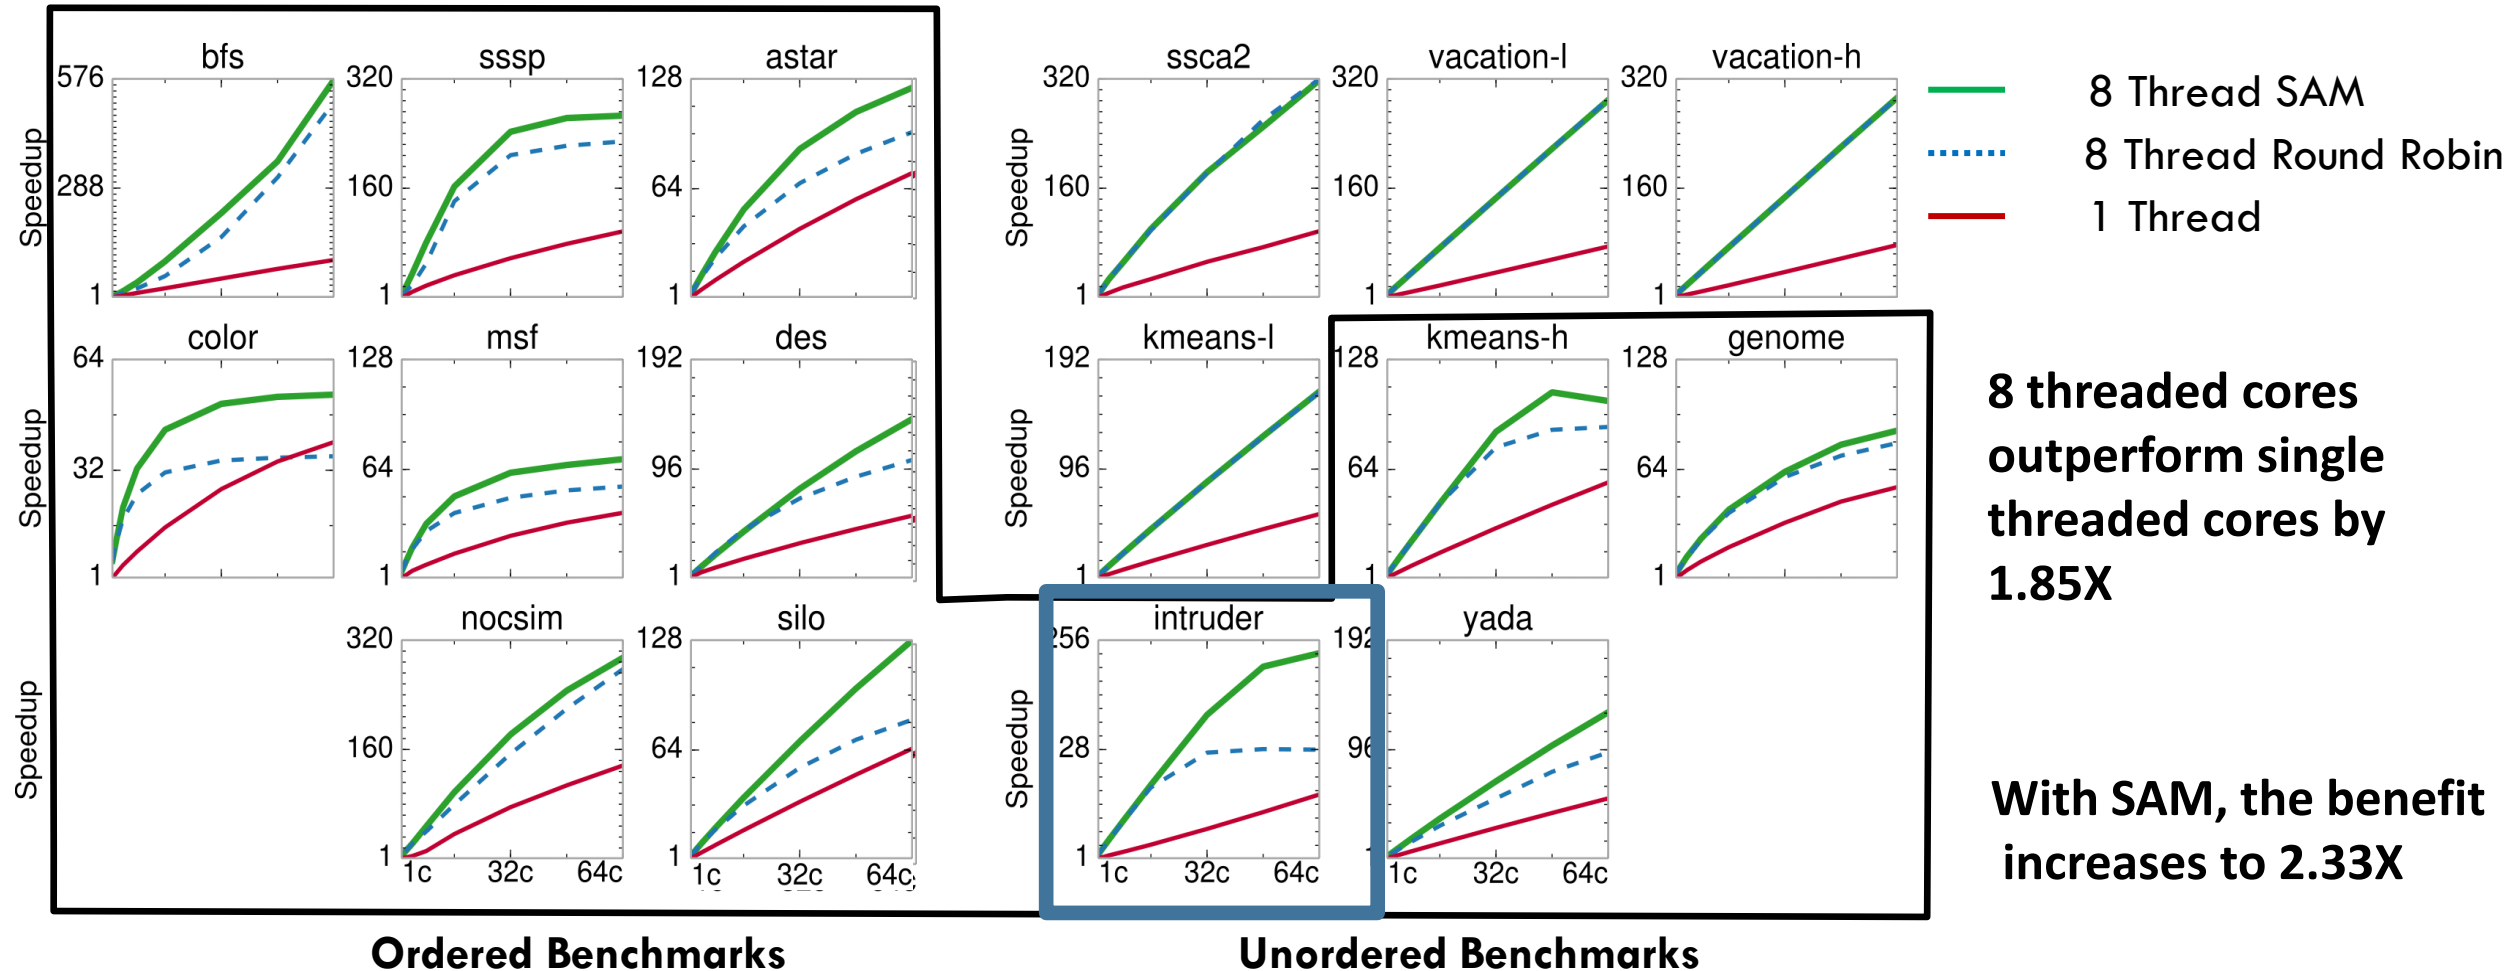
Baseline System

- Swarm + Wait-N-GoTM [Jafri et al. ASPLOS'13] conflict resolution techniques
- Cycle-accurate, event-driven, Pin-based simulator
- Model systems up to 64 cores
- Cores: 2 wide issue, up to 8 threads per core

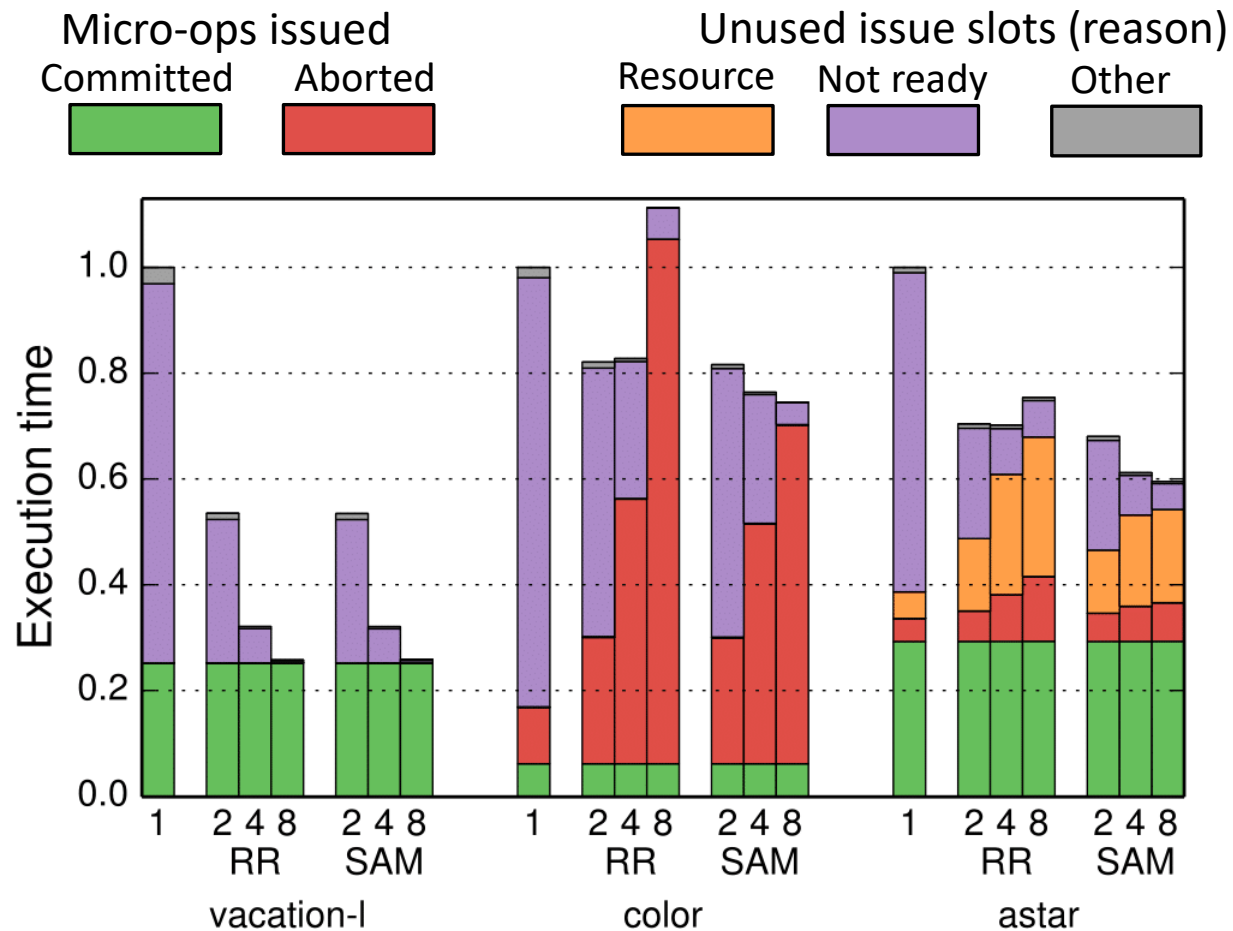
Benchmarks

- Ordered : Swarm [Jeffrey et al. MICRO'15, MICRO'16] – 8 benchmarks
- Unordered : STAMP [Minh et al. IISWC' 08] – 8 benchmarks

SAM makes multithreading more effective



Why does SAM help?



SAM matches RR when there are no pathologies

SAM reduces wasted work

SAM reduces resource stalls

Outline

Background on speculative parallelism

Pitfalls of speculative parallelism with conventional multithreading

SAM on in-order cores

SAM on out-of-order cores

SAM on out-of-order cores

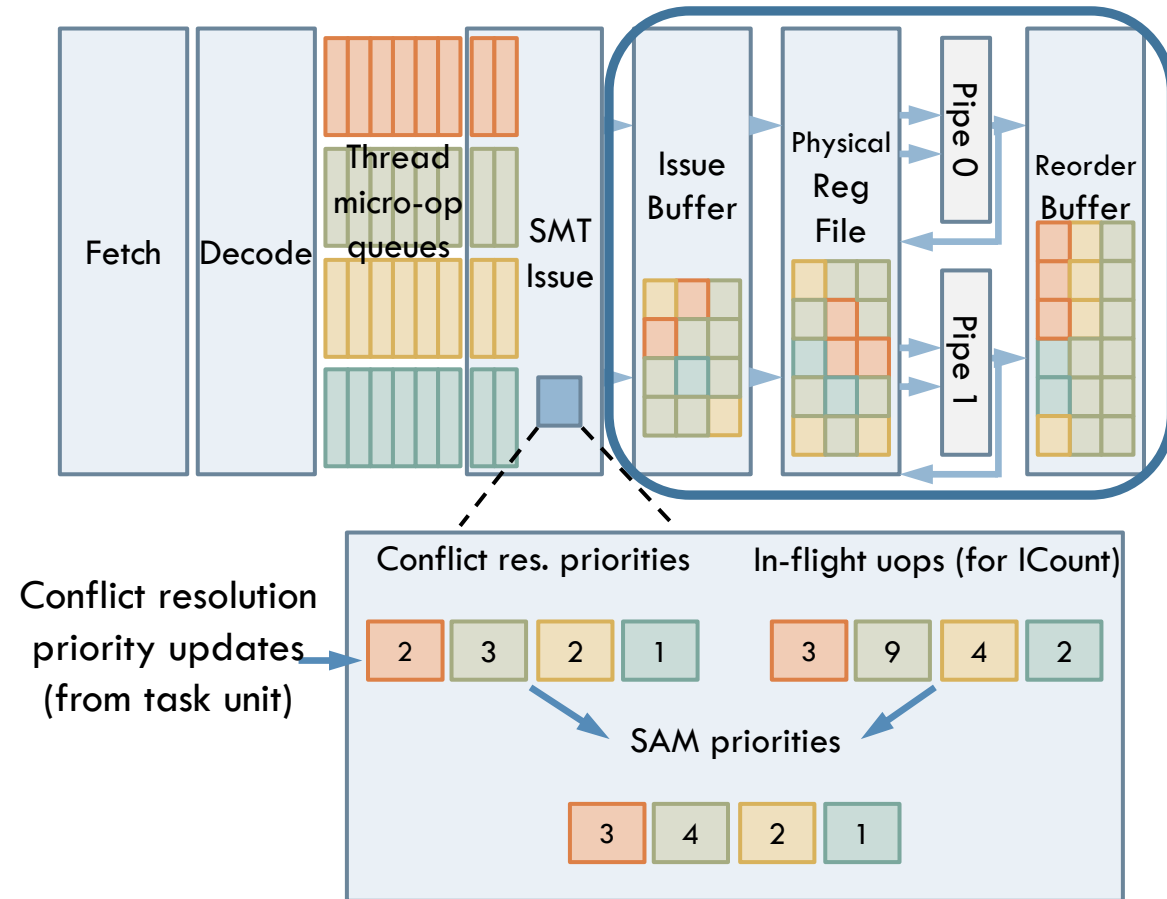
Unlike in-order cores, priorities affect pipeline efficiency

- A single thread can clog core resources
- Increased wrong path execution

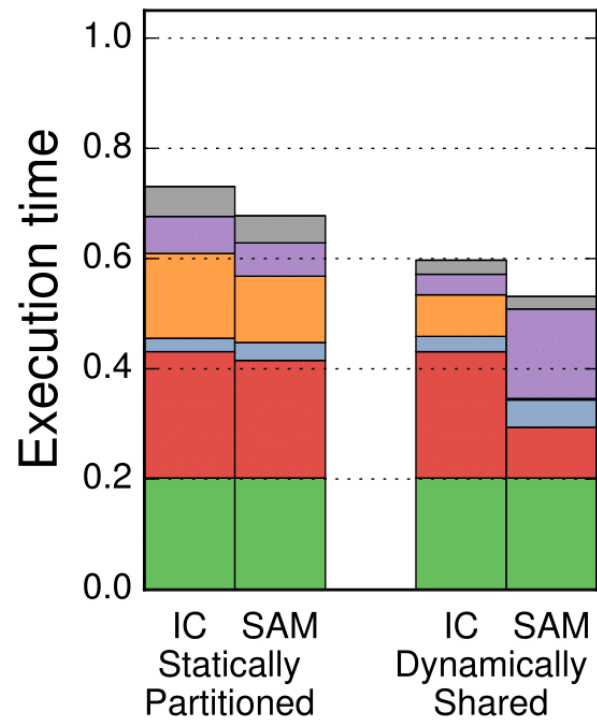
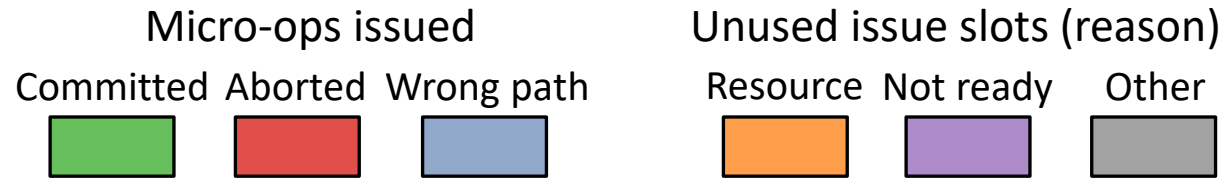
Despite these, prioritizing tasks is better

Need for aggressive prioritization affects core design

- Shared, not partitioned ROBs



SAM tradeoffs with out-of-order cores



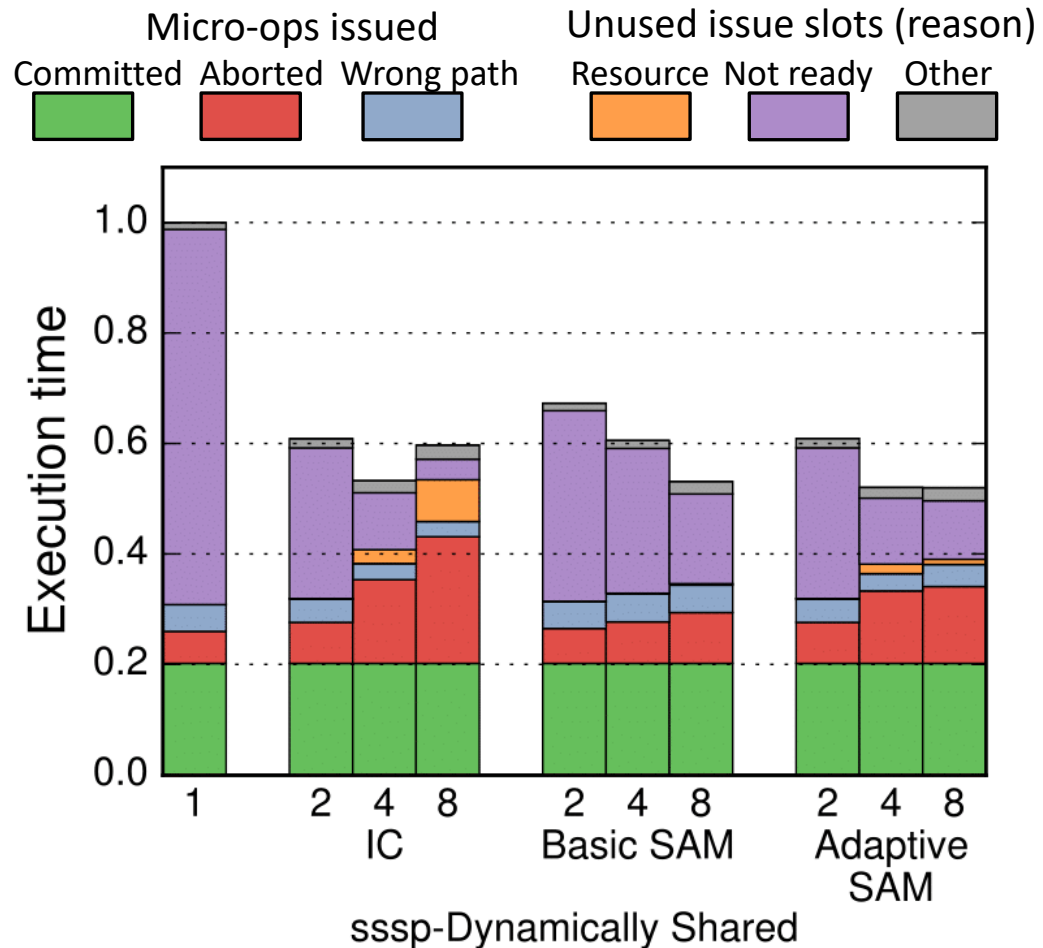
sssp – 8 threads

Baseline policy - ICount (IC)

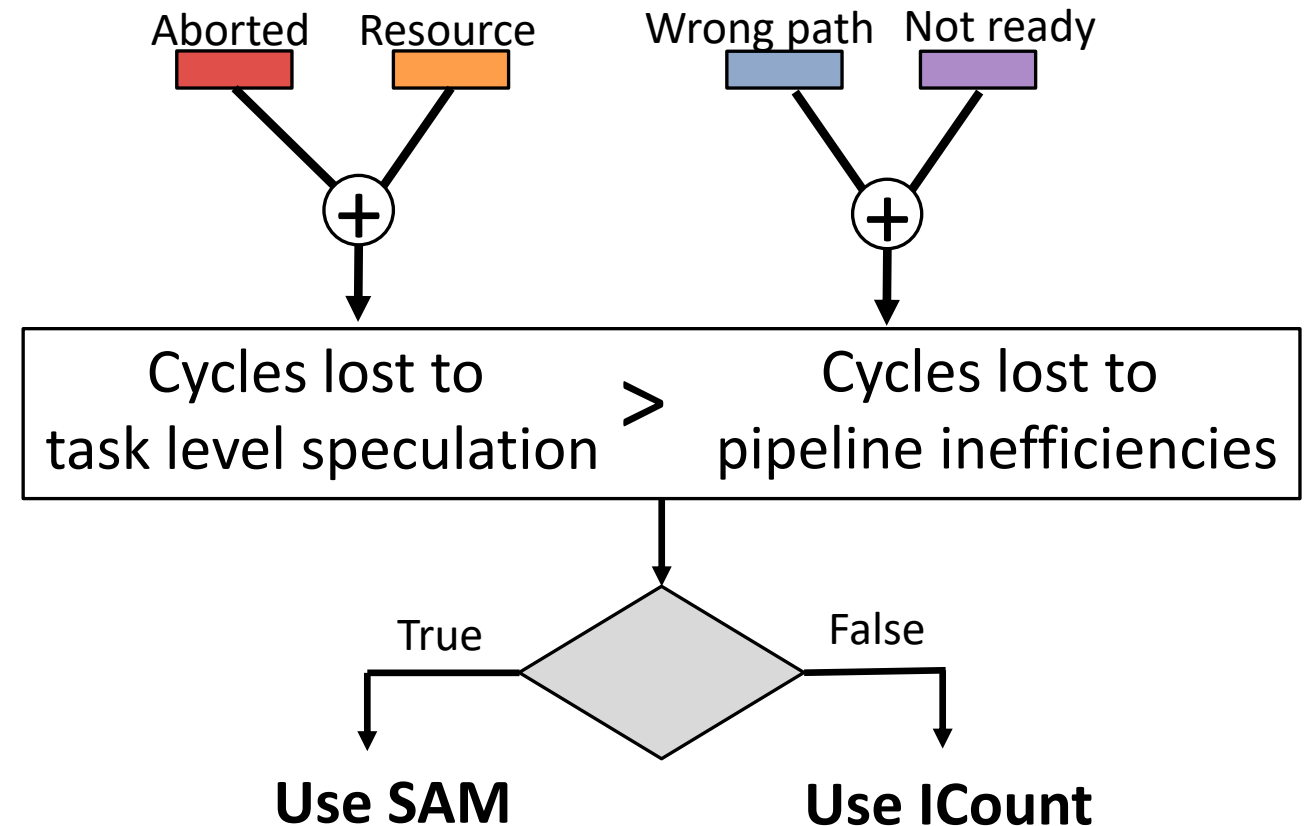
SAM is more beneficial with dynamically shared ROB
Reduces aborts + resource stalls

But reduced pipeline efficiency
Increase in wrong-path issues + not-ready stalls

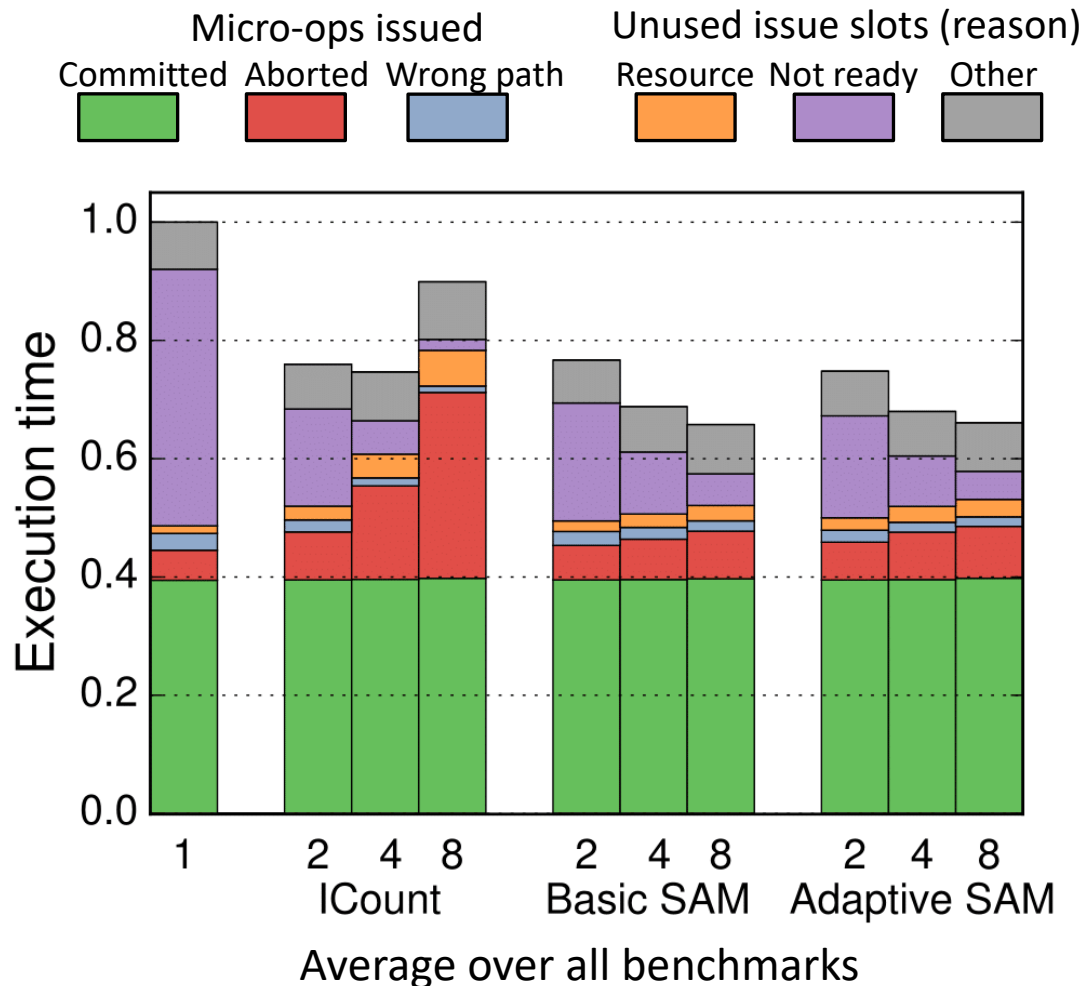
Adaptive SAM policy



Hardware counters to track cycles



SAM on OoO cores (all benchmarks)



At 8 threads / core:

- Multithreading improves performance over single threaded cores by 1.1x
- With SAM, improvement rises to 1.5x

Adaptive policy slightly increases performance at 2 and 4 threads

Conclusion

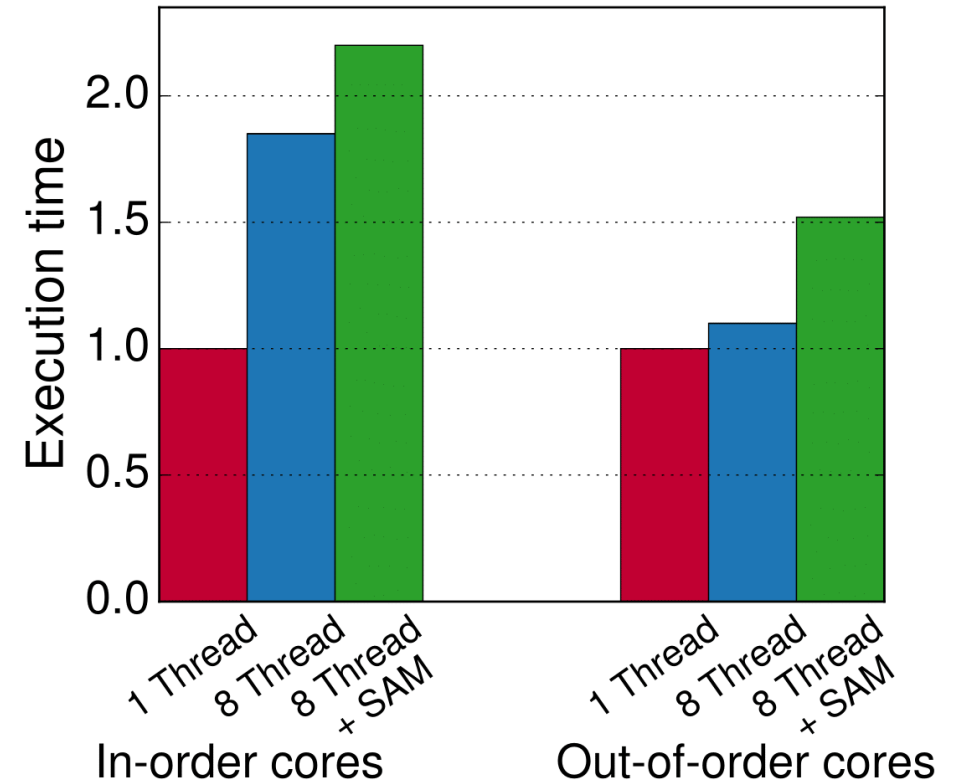
Conventional multithreading causes performance pathologies on speculative workloads

- Increase in aborted work
- Inefficient use of speculation resources

Speculation Aware Multithreading (SAM)

Prioritize threads running tasks more likely to commit

SAM makes multithreading more useful



Questions?

Conventional multithreading causes performance pathologies on speculative workloads

- Increase in aborted work
- Inefficient use of speculation resources

Speculation Aware Multithreading (SAM)

Prioritize threads running tasks more likely to commit

SAM makes multithreading more useful

